

GAN research WORK

What is a CGAN?

A Conditional Generative Adversarial Network (CGAN) is a variation of the traditional GAN that conditions the generator and discriminator on auxiliary information, such as labels or attributes.

- **Generator (G):**
 - Takes noise (zzz) and class labels (yyy).
 - Outputs synthetic images matching the desired label yyy.
 - **Discriminator (D):**
 - Takes an image and the corresponding class label yyy.
 - Determines if the image is real or fake while considering yyy.
-

3. Key Components in the Code

3.1 Generator

The **Generator** is responsible for creating synthetic images.

Key Functions:

- **nn.Embedding:**
 - Embeds class labels into dense vectors for conditioning.
 - Maps each label to a learnable vector in a high-dimensional space.
 - **torch.cat:**
 - Concatenates the noise vector (zzz) and the embedded label into a single input for the generator.
 - **nn.Tanh:**
 - Used in the last layer to scale output pixel values to $[-1, 1]$.
-

3.2 Discriminator

The **Discriminator** evaluates whether the input image is real or fake.

Key Functions:

- **nn.Embedding:**
 - Embeds class labels, similar to the generator.
- **view:**
 - Flattens input images into a 1D vector for processing.
- **torch.cat:**

- Combines the flattened image and the embedded label for conditioning.
 - **nn.Sigmoid:**
 - Outputs a probability score (real/fake) for the input.
-

4. Training Process

Steps:

1. **Train Discriminator:**
 - Compute loss for real images ($\mathcal{L}_{D_{\text{real}}}$).
 - Compute loss for fake images ($\mathcal{L}_{D_{\text{fake}}}$).
 - Update discriminator weights to minimize:

$$\mathcal{L}_D = -\log D(x, y) - \log(1 - D(G(z, y), y))$$
 2. **Train Generator:**
 - Generate fake images from z and y .
 - Compute loss based on the discriminator's output for fake images:

$$\mathcal{L}_G = -\log D(G(z, y), y)$$
 - Update generator weights to maximize \mathcal{L}_G .
-

5. Key Functions Used

5.1 Data Handling

- **transforms.ToTensor():**
 - Converts images to PyTorch tensors.
 - **transforms.Normalize():**
 - Scales image pixel values to a specified range, in this case, $[-1, 1]$.
 - **DataLoader:**
 - Handles batching and shuffling during training.
-

5.2 Model Optimization

- **nn.BCELoss():**
 - Binary cross-entropy loss used for real/fake classification.
 - Real images are labeled as 1; fake images as 0.
- **optim.Adam():**
 - Optimizer used to update the weights of the generator and discriminator.

5.3 Image Generation

- `torch.randn()`:
 - Generates random noise for the generator.
 - `torch.arange()`:
 - Creates a sequence of class labels for generating images of specific classes.
 - `vutils.save_image()`:
 - Saves generated images in a grid format for visual inspection.
-

6. Results

The model generates images conditioned on class labels. For example:

- Given label $y=3$, the generator produces an image resembling the digit "3".

Images are saved to the file `cgan_generated.png` for evaluation.

7. Applications

- **Image Synthesis**: Generate images for specific classes.
 - **Data Augmentation**: Create additional samples for imbalanced datasets.
 - **Style Transfer**: Condition on attributes like color, texture, or style.
-

8. Challenges and Improvements

- **Mode Collapse**: The generator may produce limited variations for each class.
 - Solution: Use techniques like mini-batch discrimination or Wasserstein loss.
- **Evaluation Metrics**: Use metrics like FID (Fréchet Inception Distance) for quantitative assessment.

In a **CGAN**, the generator tries to minimize the **binary cross-entropy** loss, where the discriminator is tasked with distinguishing between real and fake images.

Fréchet inception distance (FID) is a metric for quantifying the realism and diversity of images generated by generative adversarial networks (GANs).

